ISHAAN DHAWAN

Bangalore, India | 8283806083 | ishaanaadi2000@gmail.com | linkedin.com/in/ishaanaadi2000 | github.com/ishaandhawan3

SUMMARY

AI Engineer building production GenAI and applied ML systems — intelligent document processing, LLM-based extraction and reconciliation pipelines, and anomaly detection at ~97% extraction accuracy across high-volume workloads. Hands-on across the full LLM application stack: RAG, LangChain/LangGraph, Azure OpenAI, and FastAPI backends, shipping products from POC to production.

EXPERIENCE

Data Science Engineer *Jan 2026 – Present*

DIONCE Technologies, Bangalore, India

- Architected and deployed an end-to-end intelligent document processing pipeline (Azure Document Intelligence + LLM-based post-processing), achieving ~97% validated extraction accuracy on high-volume production documents and cutting manual review effort.
- Built document classification and structured extraction workflows with validation and normalization layers, exposing API-ready outputs consumed by downstream business systems.
- Designed an OCR + LLM purchase-order reconciliation system that normalizes extracted fields and auto-flags discrepancies against ERP records, reducing manual reconciliation effort.
- Architected a hybrid anomaly-detection pipeline combining domain heuristics with Isolation Forest, generating multi-factor 0–100 risk scores to prioritize suspicious transactions for review.

Data Science Intern *Jul 2025 – Jan 2026*

DIONCE Technologies, Bangalore, India

- Built a Django-based software license optimization platform for a global life-sciences client, taking it from POC to MVP and surfacing ~€500K in potential cost-saving opportunities.
- Developed backend pipelines for Excel ingestion, rule-based analytics, and AI-generated reporting, powering dashboards and exportable decision-support insights.
- Integrated Azure OpenAI into reporting workflows to generate contextual summaries and explainable business insights from analytical outputs.

AI/ML Intern *Sep 2024 – Jan 2025*

SourceFuse Technologies

- Built GenAI applications using Python, Flask, LangChain, Hugging Face, prompt engineering, and retrieval-augmented generation (RAG).
- Optimized retrieval and prompt pipelines, improving chatbot response time by 10%, retrieval efficiency by 20%, and response quality by 15%.

PROJECTS

Activity Management Backend APIs — *Python, FastAPI, AWS S3, REST APIs*

- Designed backend APIs for personalized activity generation, search, categorization, lifecycle tracking, media uploads, scoring, and feedback with secure S3-based asset handling.

Activity Engine — *Python, FastAPI, Streamlit, AWS S3*

- Built an interactive Streamlit application for personalized activity discovery and progress tracking, integrating real-time backend APIs and role-based workflows.

SKILLS

Languages: Python, SQL, C/C++

GenAI / LLM: LangChain, LangGraph, RAG, Prompt Engineering, Azure OpenAI, Azure Document Intelligence, Hugging Face

ML / DL: Scikit-learn, PyTorch, TensorFlow, Keras, NLP

Backend: FastAPI, Django, Flask, REST APIs

Data: NumPy, SciPy, PySpark, MySQL

Cloud / MLOps: Azure, GCP, AWS S3, Docker, Git, Grafana, OpenTelemetry

Tools: Streamlit, Gradio, n8n, Jupyter

EDUCATION

B.E. Computer Science Engineering — AI/ML *2021 – 2025*

Chandigarh University, Punjab, India